

Video Games ETL Pipeline


Technical & Operational Documentation

Table of Contents

- 1. General Overview
 - Business Objectives
- 2. Project Specifications
 - 2.1 Business Requirements & Management Rules
 - 2.2 Technical Data Specifications
- 3. Technical Solution & Architecture
 - 3.1 Stack and Tools
 - 3.2 Project Directory Tree
 - 3.3 Architecture Diagram
 - 1. Ingestion and Preparation (Profile: Developer)
 - 2. Manual Pipeline Execution (Profile: Developer)
 - 3. Orchestration and Deployment (Profile: Administrator)
 - 4. Data Consumption (Profile: Analysts & Decision Makers)
 - 3.4 Data Schema
 - Datalake (MongoDB NoSQL)
 - Data Warehouse (SQL)
 - 3.5 Traceability Matrix
 - 3.6 Engineering Enhancements
- 4 Deployment Guide (Administrator)
 - Step 1: Cloud Infrastructures Provisioning
 - Step 2: Codebase Cloning and Dependency Setup
 - Step 3: Launching and Initializing Airflow Services
- 5 Developer Guide
 - Direct Local ETL Execution
 - Automated Dataset Seeding
 - Database Maintenance (Date Postdating)
 - Logical Data Flow Sequencing
 - Aggregation Logic Blocks
- 6 Monitoring & Infrastructure Maintenance
 - Orchestration Engine (Airflow):
 - Data Warehouse Quality Checks:

1. General Overview

This document defines the architecture, configuration, and operation of the automated ETL pipeline. The primary objective is to extract raw customer reviews stored in a NoSQL database daily, identify community trends, and populate a relational Data Warehouse (DWH).



Business Objectives

- **Catalog Optimization:** Feature the top-rated and most appreciated games on the website's homepage and in marketing communication channels (newsletters, social media networks).
- **Data Freshness:** Maintain a day-by-day historical record of the top 15 highest-rated games based exclusively on customer reviews submitted over the last 6 months.

2. Project Specifications

2.1 Business Requirements & Management Rules

- **Rolling Window:** Strict exclusion of any product review submitted more than 6 months prior to the execution date.
- **Top 15:** Daily calculation based on average ratings and overall review volume to extract exactly the top 15 references.
- **Idempotence & Uniqueness:** Zero tolerance for duplicate entries. If the pipeline runs multiple times on the same logical day, existing records for that specific date must be overwritten and replaced (**Upsert / Replace Strategy**).

2.2 Technical Data Specifications

The raw source data originates from a compressed JSON stream ingested into MongoDB Atlas.

- **Source Stream (JSON File)**

Each entry represents a single customer review containing the following fields:

- **reviewerID:** Unique identifier of the user.
- **verified:** Boolean flag indicating if the purchase is verified (ETL filtering criterion).
- **asin:** Unique product identifier (Amazon Standard Identification Number).
- **reviewerName:** Name or handle of the author.
- **vote:** Number of helpful votes received by the review.
- **style:** Dictionary describing the product format (e.g., "Digital Download").
- **reviewText:** Text body of the review (dropped during ingestion to optimize cluster storage capacity).
- **overall:** Numerical rating ranging from 1 to 5.
- **summary:** Headline or summary title of the review.
- **unixReviewTime:** Unix Timestamp (used to calculate the dynamic 6-month rolling window).
- **reviewTime:** Formatted string date (e.g., "05 22, 2024").
- **image:** A list of URLs pointing to images uploaded by the customer.

- **Datalake (MongoDB NoSQL)**

The staging collection stores the filtered fields extracted from the raw file required by the target schema. Only the necessary fields from the file need to be fetched.

- **Target DWH (SQL)** The final target relational table features the following columns:

- Unique game identifier (ASIN)
- Calculated average rating
- Total count of users who rated the game
- Oldest review rating retained in the rolling window
- Most recent review rating recorded in the rolling window
- Calculation execution date

3. Technical Solution & Architecture

3.1 Stack and Tools

- **Programming Languages:** Python 3.11, JavaScript
- **Storage Infrastructure:** MongoDB Atlas (Source Database), PostgreSQL on Render (Data Warehouse).
- **Data Processing:** Pandas & MongoDB Aggregation Framework.
- **Orchestration Engine:** Apache Airflow (Local Environment / Server Instance).
- **IDE:** Visual Studio Code
 - **Extensions:**
 - MongoDB for VS Code
 - SQLTools
 - SQLTools PostgreSQL/Cockroach Driver
 - SQLite Viewer
- **Operating System:** Ubuntu Linux. (The Airflow Scheduler strictly requires a Unix-like OS environment to support task process handling via `os.fork`).

3.2 Project Directory Tree

```

BlentDataProject/
├── .venv_airflow/      # Dedicated virtual environment for Airflow
├── .venv_etl/          # Dedicated virtual environment for the ETL script
├── airflow_home/       # Airflow working directory (Logs, local DB)
│   ├── airflow.db      # Orchestrator SQLite database
│   └── dags/           # Airflow DAGs folder
│       └── dag_task_etl.py # Airflow 2.3 DAG definition & integrated documentation
├── doc/                # Technical and functional specifications
├── queries/            # Maintenance scripts (MongoDB/SQL migrations)
├── scripts/
│   └── run_etl.py       # Python script (Extraction, Calculations, Loading)
├── src/                # Core logic (lib_etl.py, config.py)
├── .env.template       # Secrets template (MongoDB, Postgres)
├── .gitignore          # Exclusions for virtual environments, logs, and .env
├── airflow.env         # Path environment variables (PROJECT_ROOT, AIRFLOW_HOME)
├── airflow_run_etl.sh  # Control script (Servers & execution modes)
├── README.md           # Overview and quick-start guide
├── requirements_airflow.txt # Orchestrator dependencies
└── requirements_etl.txt  # ETL script dependencies


```

3.3 Architecture Diagram

This section details data streams and architecture components mapped across different user profiles and business roles.


1. Ingestion and Preparation (Profile: Developer)

- **Objective:** Seed the Datalake with workable development datasets.
- **Process:** Loading the source JSON file using the `seed_dataLake` function, followed by a date-shifting script to map older documents into a modern 6-month execution window.
- *Ref. Diagram 1: Ingestion workflow and timestamp-shifting script*




2. Manual Pipeline Execution (Profile: Developer)

- **Objective:** Perform isolated unit testing or execute performance benchmarks.
- **Process:** Triggering the entry-point script `run_etl.py` directly from the CLI using `--scan_date` and `--platform` arguments.
- *Ref. Diagram 2: Direct Extraction, Transformation, and Loading pipeline*




3. Orchestration and Deployment (Profile: Administrator)

- **Objective:** Management of background service infrastructure and production automation.
- **Components:** A control shell script `airflow_run_etl.sh` driving 2 daemon processes (Webserver & Scheduler) paired with an `airflow.db` metadata instance.
- **Operation Modes:** Daily Schedule (Automated), Test (Unit testing), Backfill (Historical catching up).
- *Ref. Diagram 3: Airflow orchestration architecture and execution modes*



4. Data Consumption (Profile: Analysts & Decision Makers)

- **Objective:** Business trend monitoring and decision support analytics.
- **Flow:** Render PostgreSQL DWH → Standard SQL Queries → Analytics Reports / Marketing Newsletters / Web Catalog App.
- *Ref. Diagram 4: Final SQL data consumption workflow*



3.4 Data Schema

Datalake (MongoDB NoSQL)

The collection schema stores data with the following technical fields:

Column	Type	Description
asin	str	Unique product identifier (ASIN).
reviewerID	str	Unique user identifier.
overall	float	Numerical rating from 1 to 5.
verified	bool	Boolean flag indicating a verified purchase.
unixReviewTime	int	Unix Epoch Timestamp.
reviewTime	str	Formatted date text.

Data Warehouse (SQL)

The target relational table `daily_snapshot` enforces the following schema to ensure historical consistency and idempotence:

Column	Type	Description
product_id	VARCHAR(50)	Unique game identifier (ASIN).
snapshot_date	DATE	Pipeline calculation execution day (Composite Primary Key with product_id).
nb_reviews	INTEGER	Number of valid reviews aggregated within the 6-month window.
average_rating	NUMERIC(3,2)	Calculated arithmetic mean rating (1.00 to 5.00).
oldest_rating	NUMERIC(3,2)	Rating value of the oldest review in the current period.
newest_rating	NUMERIC(3,2)	Rating value of the latest review in the current period.

3.5 Traceability Matrix

Requirement Segment	Specific Requirement	Script (folder)	Implementation Status & Comments
Source Data	Raw dataset hosted in MongoDB	<code>lib_etl.py</code> (src)	Implemented. <code>connect_mongo</code> and <code>extract_and_transform</code> use <code>pymongo</code> to query the source cluster collection.
Data Warehouse	SQL Compliant Infrastructure (PostgreSQL)	<code>lib_etl.py</code> (src)	Implemented. Uses <code>sqlalchemy</code> and <code>psycopg2</code> (via explicit connection string DSN) pointing to the Render PostgreSQL instance.
Data Filtering	Restrict dataset to the last 6 months	<code>lib_etl.py</code> (src)	Implemented. <code>get_timeframe_start</code> calculates the 6-month delta, and <code>extract_and_transform</code> leverages a <code>\$match</code> operator on <code>unixReviewTime</code> .
Aggregation Rules	Extract top 15 highest-rated games	<code>lib_etl.py</code> (src)	Implemented. Aggregation pipeline enforces a <code>\$limit: top_n</code> (defaulting to 15) after sorting records by <code>average_rating</code> and <code>nb_reviews</code> .
Data Schema Matching	Product ID, Average, Count, Oldest/Newest bound ratings	<code>lib_etl.py</code> (src)	Implemented. <code>init_dwh</code> initializes the schema with appropriate constraints for <code>product_id</code> , <code>nb_reviews</code> , <code>average_rating</code> , <code>oldest_rating</code> , and <code>newest_rating</code> .
Idempotence Enforcer	Handle duplicate execution runs cleanly	<code>lib_etl.py</code> (src)	Implemented. <code>upsert_dwh</code> executes an atomic row <code>DELETE</code> for the targeted <code>snapshot_date</code> right before triggering the DataFrame <code>append</code> .
Core App Logic	Standalone Python execution script	<code>run_etl.py</code> (src)	Implemented. The primary CLI entry point script coordinates individual Extract, Transform, and Load phases sequentially.
Automation	Orchestration & Job Scheduling engine	<code>airflow_run_etl.sh</code>	Implemented. Control Bash script manages local environment properties, background daemon systems, and explicit backfilling.
Data Integrity	Filter out unverified product reviews	<code>lib_etl.py</code> (src)	Implemented. The primary MongoDB server-side aggregation stage explicitly filters records where <code>verified: True</code> .

Project Status Assessment Summary:

- **Date Management:** Transitioning legacy sample datasets (dated around 2017) to a modern execution timeline was successfully implemented using the maintenance utility `queries/datalake/change_dates.mongodb.js`. This guarantees that the dynamic 6-month historical calculations yield consistent rows during active production testing.
- **Schema Consistency:** The database creation DDL block within `src/lib_etl.py` configures a composite `PRIMARY KEY (product_id, snapshot_date)`. This fulfills the requirement to eliminate multi-run data pollution while preserving true daily metrics tracking for individual games over time.

3.6 Engineering Enhancements

1. **Performance Optimization:** Offloading computations using the native MongoDB `aggregate` pipeline reduces network overhead and minimizes data transfer into Python.
2. **Pipeline Robustness:** Implementing contextual SQLAlchemy database transactions (`db_dwh.begin()`) protects data warehouse integrity during append sequences.
3. **Security Architecture:** Completely decoupled secret configurations via local `.env` environment isolation.

4 Deployment Guide (Administrator)

Step 1: Cloud Infrastructures Provisioning

1. MongoDB Atlas (Source Datalake):

- Deploy a Shared/Free M0 tier cluster instance.
- Under **Network Access**, whitelist your runner IP address (or `0.0.0.0/0` during development).
- Under **Database Access**, create a dedicated application user account with `readWrite` access privileges (required to seed initial testing objects).
- Retrieve your secure cluster connection URI string (`mongodb+srv://...`).

2. Render PostgreSQL (Data Warehouse Target):

- Provision a new managed **PostgreSQL** cloud instance.
- Configure database user credentials.
- Copy down the **External Connection String**.

Step 2: Codebase Cloning and Dependency Setup

```
# 1. Clone the project repository
git clone https://github.com/votre-compte/BlentDataProject.git
cd BlentDataProject

# 2. Configure application secrets
cp .env.template .env
# Edit the newly created .env file with your custom MongoDB and PostgreSQL strings

# 3. Initialize separate virtual environments
python3.11 -m venv .venv_airflow
python3.13 -m venv .venv_etl

# 4. Install requirements packages across environments
source .venv_airflow/bin/activate && pip install -r requirements_airflow.txt && deactivate
source .venv_etl/bin/activate && pip install -r requirements_etl.txt && deactivate
```

Step 3: Launching and Initializing Airflow Services

Execute the primary orchestration control script to spin up background operations:

```
chmod +x airflow_run_etl.sh
./airflow_run_etl.sh
```

This management script automatically orchestrates:

1. The provisioning of the `airflow_home` directory structures and initializing `airflow.db`.
2. The programmatic injection of an `admin` security profile account.
3. The clean initialization of the Airflow **Scheduler** and **Webserver** (Port 8080) as background daemons.

5 Developer Guide

Direct Local ETL Execution

Developers can bypass the Airflow UI layers to run unit tests on the pipeline processing layer directly:

```
source .venv_etl/bin/activate

# Execute the pipeline using the current runtime date context
python scripts/run_etl.py

# Execute the pipeline passing an explicit target date (ISO format string)
python scripts/run_etl.py --scan_date 2024-05-20 --platform Terminal
```

Automated Dataset Seeding

The entry point script `run_etl.py` handles bootstrap actions during its very first initialization sequence:

- **MongoDB Atlas Seeding:** If the destination staging database contains zero documents, the system triggers `seed_datalake` to automatically parse and load items from the packaged JSON file.
- **PostgreSQL DDL Initialization:** The inner utility `init_dwh` automatically provisions the target relational schema layout `daily_snapshot` alongside optimized tracking indices if they do not exist.

Database Maintenance (Date Postdating)

If testing datasets fall outside your 6-month processing calculation scope:

- Open your workspace using Visual Studio Code.
- Access your target server connection through the official MongoDB Extension.
- Open the query scratchpad file located at `queries/datalake/change_dates.mongodb.js`.
- Click the interactive execution **Play** icon to shift timestamps forward.

Logical Data Flow Sequencing

1. **Extract:** Execute a targeted aggregation pipeline matching documents where `unixReviewTime` (Date - 6 months) while verifying that `verified: true`.
2. **Transform:**
 - Sort records chronologically.
 - Group values under distinct `asin` document clusters.
 - Compute arithmetic rating means and isolate initial boundaries (`$first`, `$last`).
 - Order structural output by `average_rating` and `nb_reviews` fields descending.
 - Enforce a processing limit constraint to truncate results down to the top 15 records.
3. **Load:**
 - Open an isolated database transaction window.
 - Clear out prior target tracking entries matching the current operational `snapshot_date`.
 - Perform a high-speed relational bulk insert using the processed data framework.

Aggregation Logic Blocks

The core metric computation engine utilizes the server-side MongoDB pipeline within `extract_and_transform`. This architecture offloads calculations, filtering, and statistical sorting tasks to the hosting database machine.

```
# Extracted pipeline code block definition
pipeline = [
    {"$match": {"unixReviewTime": {"$gte": timeframe_start}, "verified": True}},
    {"$sort": {"unixReviewTime": 1}},
    {"$group": {
        "_id": "$asin",
        "nb_reviews": {"$sum": 1},
        "average_rating": {"$avg": "$overall"},
        "oldest_rating": {"$first": "$overall"},
        "newest_rating": {"$last": "$overall"}
    }},
    {"$sort": {"average_rating": -1, "nb_reviews": -1}},
    {"$limit": 15}
]
```

6 Monitoring & Infrastructure Maintenance

Orchestration Engine (Airflow):

- **Job Triggering & Service CLI Interfaces:**
 - **Spin Up Core Servers:** Run `./airflow_run_etl.sh` without passing any command-line parameters.
 - **Isolated Date Unit Verification:** Run the shell utility passing a target ISO date string parameter (e.g., `./airflow_run_etl.sh 2024-05-20`). This boots an internal `airflow tasks test` routine to validate pipeline logic without recording state mutations in the scheduler history database.
 - **Historical Backfilling:** To process an extensive chronological window on demand, invoke:
`./airflow_run_etl.sh <start_date> <end_date>`.
 - *Important Note:* The global absolute lower-bound start constraints are defined via configuration properties within `airflow.env`. Verify that valid source records exist on your MongoDB instances before backfilling target periods.
- **System Health Monitoring:**
 - **Web Management Portal:** Connect to the running Airflow server UI by loading `http://localhost:8080`.
 - **Control Panel Dashboard:** Toggle the switch to activate the target `daily_scan` scheduling graph.

Data Warehouse Quality Checks:

Run verification scripts directly against your tables to check the computed results for a specific execution run:

```
SELECT * FROM public.daily_snapshot
WHERE snapshot_date = <TARGET_DATE YYYY-MM-DD>;
```

J. Vallée - 2026-05-20